

project

loading libraries

```
library(readr)
```

```
## Warning: package 'readr' was built under R version 3.6.3
```

```
library(DataExplorer)
```

```
## Warning: package 'DataExplorer' was built under R version 3.6.3
```

```
library(caret)
```

```
## Warning: package 'caret' was built under R version 3.6.3
```

```
## Loading required package: lattice
```

```
## Loading required package: ggplot2
```

```
## Warning: package 'ggplot2' was built under R version 3.6.3
```

```
library(xgboost)
```

```
## Warning: package 'xgboost' was built under R version 3.6.3
```

```
library(ISLR)
```

```
## Warning: package 'ISLR' was built under R version 3.6.3
```

```
library(caTools)
```

```
## Warning: package 'caTools' was built under R version 3.6.3
```

```
library(summarytools)
```

```
## Warning: package 'summarytools' was built under R version 3.6.3
```

```
## Registered S3 method overwritten by 'pryr':
```

```
##   method      from
```

```
##   print.bytes Rcpp
```

```
## For best results, restart R session and update pander using devtools:: or remotes::install_github('r')
```

```
library(DMwR)
```

```
## Warning: package 'DMwR' was built under R version 3.6.3
```

```
## Loading required package: grid
```

```
## Registered S3 method overwritten by 'quantmod':  
##   method             from  
##   as.zoo.data.frame zoo
```

```
library(rattle)
```

```
## Warning: package 'rattle' was built under R version 3.6.3
```

```
## Loading required package: tibble
```

```
## Warning: package 'tibble' was built under R version 3.6.3
```

```
##  
## Attaching package: 'tibble'
```

```
## The following object is masked from 'package:summarytools':  
##  
##   view
```

```
## Loading required package: bitops
```

```
## Rattle: A free graphical interface for data science with R.  
## Version 5.4.0 Copyright (c) 2006-2020 Togaware Pty Ltd.  
## Type 'rattle()' to shake, rattle, and roll your data.
```

```
##  
## Attaching package: 'rattle'
```

```
## The following object is masked from 'package:xgboost':  
##  
##   xgboost
```

load dataset:

```
Cars_data <- read_csv("C:/Users/daoud/Downloads/PGP DSBA/productive modeling/week 5 project/Cars-dataset.csv")
```

```
## Parsed with column specification:  
## cols(  
##   Age = col_double(),  
##   Gender = col_character(),
```

```
## Engineer = col_double(),
## MBA = col_double(),
## `Work Exp` = col_double(),
## Salary = col_double(),
## Distance = col_double(),
## license = col_double(),
## Transport = col_character()
## )
```

```
#View(Cars_data)
```

Exploratory Data Analysis

```
summarytools::view(dfSummary(Cars_data))
```

```
## Switching method to 'browser'
```

```
## Output file written: C:\Users\daoud\AppData\Local\Temp\RtmpCUAqn1\file3118d11037.html
```

Observation : 1- we have 418 employee with 9 variable ,Gender and Transport are character , the other variable are numeric . 2- we have only one missing value : MBA . 3- column name “Work Exp” will change to “Work_Exp” . 4- 19.9% of employee use “2Wheeler”, 8.4% use “Car”, 71.8% use “Public Transport”. 5- dependent variable is “Transport” , independent variable are : “Age”, “Gender”, “Engineer”, “MBA”, “Work Exp”, “Salary”, “Distance”, “license” ## challenging problem : we have 3 classes on a target variable , it should be 2 only. the task was to predict whether or not an employee will use Car as a mode of transport. there are two methods to solve the problem : “levels” or “ifelse”, for today we will use “ifelse” to assign 1 for “Car” and 0 for others " Public_Transport , 2Wheeler " as Transport_car .

```
summary(Cars_data)
```

```
##      Age      Gender      Engineer      MBA
## Min.   :18.00  Length:418    Min.    :0.0000  Min.    :0.0000
## 1st Qu.:25.00  Class :character 1st Qu.:0.2500 1st Qu.:0.0000
## Median :27.00  Mode  :character  Median :1.0000 Median :0.0000
## Mean   :27.33                Mean   :0.7488 Mean   :0.2614
## 3rd Qu.:29.00                3rd Qu.:1.0000 3rd Qu.:1.0000
## Max.   :43.00                Max.    :1.0000 Max.    :1.0000
##                                     NA's    :1
##      Work Exp      Salary      Distance      license
## Min.   : 0.000  Min.    : 6.500  Min.    : 3.20  Min.    :0.0000
## 1st Qu.: 3.000  1st Qu.: 9.625  1st Qu.: 8.60  1st Qu.:0.0000
## Median : 5.000  Median :13.000  Median :10.90  Median :0.0000
## Mean   : 5.873  Mean   :15.418  Mean   :11.29  Mean   :0.2033
## 3rd Qu.: 8.000  3rd Qu.:14.900  3rd Qu.:13.57  3rd Qu.:0.0000
## Max.   :24.000  Max.   :57.000  Max.    :23.40  Max.    :1.0000
##
##      Transport
## Length:418
## Class :character
## Mode  :character
```

```
##
##
##
##
```

```
str(Cars_data)
```

```
## tibble [418 x 9] (S3: spec_tbl_df/tbl_df/tbl/data.frame)
## $ Age      : num [1:418] 28 24 27 25 25 21 23 23 24 28 ...
## $ Gender   : chr [1:418] "Male" "Male" "Female" "Male" ...
## $ Engineer : num [1:418] 1 1 1 0 0 0 1 0 1 1 ...
## $ MBA      : num [1:418] 0 0 0 0 0 0 1 0 0 0 ...
## $ Work Exp : num [1:418] 5 6 9 1 3 3 3 0 4 6 ...
## $ Salary   : num [1:418] 14.4 10.6 15.5 7.6 9.6 9.5 11.7 6.5 8.5 13.7 ...
## $ Distance : num [1:418] 5.1 6.1 6.1 6.3 6.7 7.1 7.2 7.3 7.5 7.5 ...
## $ license  : num [1:418] 0 0 0 0 0 0 0 0 0 1 ...
## $ Transport: chr [1:418] "2Wheeler" "2Wheeler" "2Wheeler" "2Wheeler" ...
## - attr(*, "spec")=
## .. cols(
## ..   Age = col_double(),
## ..   Gender = col_character(),
## ..   Engineer = col_double(),
## ..   MBA = col_double(),
## ..   `Work Exp` = col_double(),
## ..   Salary = col_double(),
## ..   Distance = col_double(),
## ..   license = col_double(),
## ..   Transport = col_character()
## .. )
```

```
Cars_data <- na.omit(Cars_data) # drop missing value
names(Cars_data)[names(Cars_data)=="Work Exp"]<-"Work_Exp" # change name without space .
Cars_data$Transport_car <- ifelse(Cars_data$Transport=="Car",1,0) # convert 'Car' to 1 else to 0
Cars_data$Gender <- ifelse(Cars_data$Gender == "Male",1,0) # convert 'Male' to 1 else to 0
prop.table(table(Cars_data$Transport))
```

```
##
##      2Wheeler      Car Public Transport
##      0.19904077      0.08393285      0.71702638
```

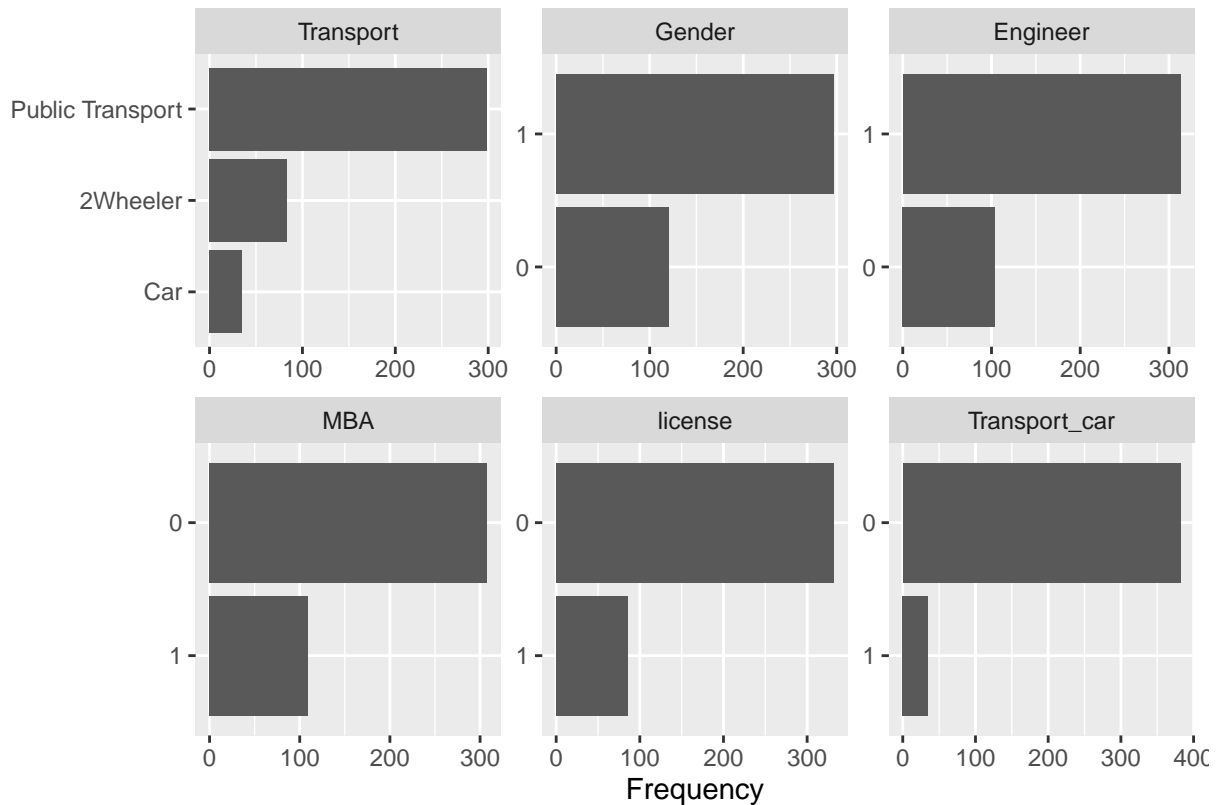
```
Cars_data
```

```
## # A tibble: 417 x 10
##   Age Gender Engineer  MBA Work_Exp Salary Distance license Transport
##   <dbl> <dbl>   <dbl> <dbl>   <dbl> <dbl>   <dbl>   <dbl> <chr>
## 1  28     1       1     0       5  14.4     5.1     0 2Wheeler
## 2  24     1       1     0       6  10.6     6.1     0 2Wheeler
## 3  27     0       1     0       9  15.5     6.1     0 2Wheeler
## 4  25     1       0     0       1   7.6     6.3     0 2Wheeler
## 5  25     0       0     0       3   9.6     6.7     0 2Wheeler
## 6  21     1       0     0       3   9.5     7.1     0 2Wheeler
## 7  23     1       1     1       3  11.7     7.2     0 2Wheeler
```


```
## 8      23      1      0      0      0      6.5      7.3      0 2Wheeler
## 9      24      1      1      0      4      8.5      7.5      0 2Wheeler
## 10     28      1      1      0      6     13.7      7.5      1 2Wheeler
## # ... with 407 more rows, and 1 more variable: Transport_car <dbl>
```

Observation : 1- there is one missing value and we drop it. 2- we have 19.9% use 2Wheeler and 8.39% use car and 71.7% use Public Transport. ## normality distribution : # visualization:

```
plot_bar(Cars_data)
```




```
plot_histogram(Cars_data)
```




observation : 1- most of employee use Public Transport. 2- most of employee are Male. 3- most of employee have Engineer Degree. 4- most of employee have MBA Degree. 5- most of employee don't have License. 6- both Age and Distance are normally distributed. 7- both Salary and Work_Exp are skewed right, possible outlier. # variables relationship :

```
ggplot(data = Cars_data, mapping = aes(x = Transport, y = Salary)) +  
  geom_boxplot()
```





observation : high Salary employees they use Car for Trasport.

```
ggplot(data = Cars_data) +  
  geom_point(mapping = aes(x = Salary, y = Work_Exp))
```




observation: 1- there is linear relationship between Salary and Work_Exp , the higher Work_Exp the higher Salary . 2- there are

```
par(mfrow=c(1,2))
hist(Cars_data$Transport_car[Cars_data$Gender==1],col = "blue",xlab = "Gender_Male",main = "Histogram T
hist(Cars_data$Transport_car[Cars_data$Gender==0],col = "blue",xlab = "Gender_Female",main = "Histogram
```





observation: most of Male and Female employee don't use car , but Male employee use car for those how use Car as Transport.

```
par(mfrow=c(1,2))
boxplot(Cars_data$Work_Exp[Cars_data$Transport_car==1],col = "blue",xlab = "car Transport",main = "Hist
boxplot(Cars_data$Work_Exp[Cars_data$Transport_car==0],col = "blue",xlab = "without car Transport",main
```




car Transport




without car Transport

observation: most of employees that have more then 15 years of Work Experience use Car for transport.

```
plot(density(Cars_data$Work_Exp),main="salary")
```





```
par(mfrow=c(1,2))
boxplot(Cars_data$Salary[Cars_data$Transport_car==1],col = "blue",xlab = "car Transport",main = "Histogram of Salary for car Transport")
boxplot(Cars_data$Salary[Cars_data$Transport_car==0],col = "blue",xlab = "without car Transport",main = "Histogram of Salary without car Transport")
```



observation: most of employees that have more then 30K of Salary use Car for transport.

```
par(mfrow=c(1,2))
hist(Cars_data$Transport_car[Cars_data$license==1],col = "blue",xlab = "with license",main = "Histogram")
hist(Cars_data$Transport_car[Cars_data$license==0],col = "blue",xlab = "without license",main = "Histogram")
```




```
sum(Cars_data$Transport_car==1 & Cars_data$license==0) # to find how many have Car without license
```

```
## [1] 6
```

observation: 1- most of employees don't have license don't have Car. 2- there are 6 employees have Car but don't have License, maybe they have personal driver.


```
par(mfrow=c(1,2))
boxplot(Cars_data$Distance[Cars_data$Transport_car==1],col = "blue",xlab = "with car",main = "Histogram with car")
boxplot(Cars_data$Distance[Cars_data$Transport_car==0],col = "blue",xlab = "without car",main = "Histogram without car")
```

Histogram Distance



with car

Histogram Distance




without car

observation: 1- most employees how live 15 KM or more from office have Car for Transport. 2- all employees live less then 20 KM don't have Car.

correlation:

```
# we drop Transport for now .  
Cars_data1<- Cars_data[, -c(9)]  
plot_correlation(Cars_data1)
```



observation: 1- as we found from the graph hight correlation between Work_exp and Salary 0.93 . 2- 0.86 correlation between Salary and Age. 3- agian hight correlation between Salary and Transport_car 0.81. # split Dataset to Train and Test :

```
set.seed(199)
Cars_data1<- as.data.frame(Cars_data1)
Cars_data1$Transport_car <- ifelse(Cars_data1$Transport==1,"Yes","No")

sample = sample.split(Cars_data1,SplitRatio = 0.75) # 75% train data , 25% test data
training = subset(Cars_data1,sample == TRUE)
testing = subset(Cars_data1,sample == FALSE)
nrow(training)
```

```
## [1] 278
```

```
nrow(testing)
```

```
## [1] 139
```

```
# the data split is equal between Train and Test with original dataset.
prop.table(table(Cars_data1$Transport_car))
```

```
##
##      No      Yes
## 0.91606715 0.08393285
```

```
prop.table(table(training$Transport_car))
```

```
##  
##           No           Yes  
## 0.91726619 0.08273381
```

```
prop.table(table(testing$Transport_car))
```

```
##  
##           No           Yes  
## 0.91366906 0.08633094
```

```
training$Transport_car<-as.factor(training$Transport_car) # to be Factor  
testing$Transport_car<-as.factor(testing$Transport_car)   # to be Factor
```

Modeling :

Setting up the general parameters for training multiple models:

```
set.seed(213)  
Crul<- trainControl(  
  method = "repeatedcv",  
  number = 5,      # number of folds  
  repeats = 10,    # repeated k-fold cross-validation  
  p = 10,  
  allowParallel = TRUE,  
  classProbs = TRUE,  
  summaryFunction = twoClassSummary  
)
```

rpart: Single decision tree :

```
rpart_model <- train(Transport_car ~ ., data = training,  
  method = "rpart",  
  minbucket = 10,  
  cp = 0,  
  tuneLength = 10,  
  trControl = Crul)
```

```
## Warning in train.default(x, y, weights = w, ...): The metric "Accuracy" was not  
## in the result set. ROC will be used instead.
```

```
rpart_model
```




```
## CART
##
## 278 samples
## 8 predictor
## 2 classes: 'No', 'Yes'
##
## No pre-processing
## Resampling: Cross-Validated (5 fold, repeated 10 times)
## Summary of sample sizes: 222, 222, 223, 223, 222, 222, ...
## Resampling results across tuning parameters:
##
##   cp          ROC        Sens       Spec
## 0.0000000 0.8657255 0.9843137 0.720
## 0.0821256 0.8657941 0.9866667 0.741
## 0.1642512 0.8657941 0.9866667 0.741
## 0.2463768 0.8656373 0.9862745 0.745
## 0.3285024 0.8656373 0.9862745 0.745
## 0.4106280 0.8656373 0.9862745 0.745
## 0.4927536 0.8656373 0.9862745 0.745
## 0.5748792 0.8656373 0.9862745 0.745
## 0.6570048 0.8558333 0.9866667 0.725
## 0.7391304 0.6127549 0.9945098 0.231
##
## ROC was used to select the optimal model using the largest value.
## The final value used for the model was cp = 0.1642512.
```

```
varimp<-varImp(rpart_model)
print(varimp)
```

```
## rpart variable importance
##
##           Overall
## Salary      100.00
## Work_Exp    79.06
## Age         75.99
## Distance    65.22
## license     28.84
## MBA          0.00
## Engineer     0.00
## Gender       0.00
```

```
fancyRpartPlot(rpart_model$finalModel)
```



Rattle 2020-Jul-03 23:31:19 daoud

obsevation: 1- Variable Importance : salary is the most important variable 2- Work_Exp, Age, Distance, license, by sort. 3- 91% employee without Car with have Salary<38 and Distance<18 .

#Accuracy:

```
rpart_pred_test <- predict(rpart_model, newdata =testing[,1:8], type = "raw")
caret::confusionMatrix(rpart_pred_test, testing$Transport_car)
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction  No  Yes
##           No 127  3
##           Yes  0  9
##
##           Accuracy : 0.9784
##           95% CI : (0.9382, 0.9955)
##           No Information Rate : 0.9137
##           P-Value [Acc > NIR] : 0.001664
##
##           Kappa : 0.8457
##
##           McNemar's Test P-Value : 0.248213
##
##           Sensitivity : 1.0000
##           Specificity : 0.7500
##           Pos Pred Value : 0.9769
```

```
##          Neg Pred Value : 1.0000
##          Prevalence : 0.9137
##          Detection Rate : 0.9137
##          Detection Prevalence : 0.9353
##          Balanced Accuracy : 0.8750
##
##          'Positive' Class : No
##
```

KNN:

```
knn<-train(
  Transport_car~.,
  data = training,
  method="knn",
  #preProcess = c("center", "scale"),
  tuneLength = 3,
  trControl = Crul)
```

```
## Warning in train.default(x, y, weights = w, ...): The metric "Accuracy" was not
## in the result set. ROC will be used instead.
```

#Accuracy:

```
knn_pred_test <- predict(knn, newdata =testing[,1:8], type = "raw")
caret::confusionMatrix(knn_pred_test, testing$Transport_car)
```

```
## Confusion Matrix and Statistics
##
##          Reference
## Prediction  No Yes
##          No 127  2
##          Yes  0 10
##
##          Accuracy : 0.9856
##          95% CI : (0.949, 0.9983)
##          No Information Rate : 0.9137
##          P-Value [Acc > NIR] : 0.0003537
##
##          Kappa : 0.9013
##
##          Mcnemar's Test P-Value : 0.4795001
##
##          Sensitivity : 1.0000
##          Specificity : 0.8333
##          Pos Pred Value : 0.9845
##          Neg Pred Value : 1.0000
##          Prevalence : 0.9137
##          Detection Rate : 0.9137
##          Detection Prevalence : 0.9281
```

```
##          Balanced Accuracy : 0.9167
##
##          'Positive' Class : No
##
```

model naive base:

```
naive_base<-train(
  Transport_car~.,
  data = training,
  method="naive_bayes",
  trControl = Crul)
```

```
## Warning in train.default(x, y, weights = w, ...): The metric "Accuracy" was not
## in the result set. ROC will be used instead.
```

```
#Accuracy:
```

```
naive_pred_test <- predict(naive_base, newdata =testing[,1:8], type = "raw")
caret::confusionMatrix(naive_pred_test, testing$Transport_car)
```

```
## Confusion Matrix and Statistics
##
##          Reference
## Prediction  No  Yes
##          No 127   1
##          Yes  0  11
##
##          Accuracy : 0.9928
##          95% CI : (0.9606, 0.9998)
##          No Information Rate : 0.9137
##          P-Value [Acc > NIR] : 5.011e-05
##
##          Kappa : 0.9526
##
##          Mcnemar's Test P-Value : 1
##
##          Sensitivity : 1.0000
##          Specificity : 0.9167
##          Pos Pred Value : 0.9922
##          Neg Pred Value : 1.0000
##          Prevalence : 0.9137
##          Detection Rate : 0.9137
##          Detection Prevalence : 0.9209
##          Balanced Accuracy : 0.9583
##
##          'Positive' Class : No
##
```

Logistic Regression :

```
glm<-train(  
  Transport_car~.,  
  data = training,  
  method = "glm",  
  family = "binomial",  
  trControl = Crul  
)
```

```
## Warning in train.default(x, y, weights = w, ...): The metric "Accuracy" was not  
## in the result set. ROC will be used instead.
```

```
## Warning: glm.fit: algorithm did not converge
```

```
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
```

```
## Warning: glm.fit: algorithm did not converge
```

```
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
```

```
## Warning: glm.fit: algorithm did not converge
```

```
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
```

```
## Warning: glm.fit: algorithm did not converge
```

```
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
```

```
## Warning: glm.fit: algorithm did not converge
```

```
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
```

```
## Warning: glm.fit: algorithm did not converge
```

```
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
```

```
## Warning: glm.fit: algorithm did not converge
```

```
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
```

```
## Warning: glm.fit: algorithm did not converge
```

```
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
```

```
## Warning: glm.fit: algorithm did not converge
```

```
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
```

```
## Warning: glm.fit: algorithm did not converge

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: algorithm did not converge

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: algorithm did not converge

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: algorithm did not converge

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: algorithm did not converge

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: algorithm did not converge

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: algorithm did not converge

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: algorithm did not converge

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: algorithm did not converge

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: algorithm did not converge

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: algorithm did not converge

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
```

Warning: glm.fit: algorithm did not converge

Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

Warning: glm.fit: algorithm did not converge

Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

Warning: glm.fit: algorithm did not converge

Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

Warning: glm.fit: algorithm did not converge

Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

Warning: glm.fit: algorithm did not converge

Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

Warning: glm.fit: algorithm did not converge

Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

Warning: glm.fit: algorithm did not converge

Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

Warning: glm.fit: algorithm did not converge

Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

Warning: glm.fit: algorithm did not converge

Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

Warning: glm.fit: algorithm did not converge

Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

Warning: glm.fit: algorithm did not converge

Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

Warning: glm.fit: algorithm did not converge

Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

```
## Warning: glm.fit: algorithm did not converge

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: algorithm did not converge

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: algorithm did not converge

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: algorithm did not converge

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: algorithm did not converge

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: algorithm did not converge

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: algorithm did not converge

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: algorithm did not converge

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: algorithm did not converge

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: algorithm did not converge

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: algorithm did not converge

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
```



```
## Warning: glm.fit: algorithm did not converge
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
## Warning: glm.fit: algorithm did not converge
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
## Warning: glm.fit: algorithm did not converge
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
## Warning: glm.fit: algorithm did not converge
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
## Warning: glm.fit: algorithm did not converge
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
## Warning: glm.fit: algorithm did not converge
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
#Accuracy:
```

```
glm_pred_test <- predict(glm, newdata =testing[,1:8], type = "raw")
caret::confusionMatrix(glm_pred_test, testing$Transport_car)
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction  No Yes
##           No 126  1
##           Yes  1 11
##
##           Accuracy : 0.9856
##           95% CI : (0.949, 0.9983)
##           No Information Rate : 0.9137
##           P-Value [Acc > NIR] : 0.0003537
##
##           Kappa : 0.9088
##
## Mcnemar's Test P-Value : 1.0000000
##
##           Sensitivity : 0.9921
##           Specificity : 0.9167
##           Pos Pred Value : 0.9921
##           Neg Pred Value : 0.9167
##           Prevalence : 0.9137
##           Detection Rate : 0.9065
##           Detection Prevalence : 0.9137
##           Balanced Accuracy : 0.9544
##
##           'Positive' Class : No
##
```

Random Forest :

```
rf<-train(  
  Transport_car~.,  
  data = training,  
  method = "rf",  
  ntree = 50,  
  maxdepth = 7,  
  tuneLength = 20,  
  trControl = Crul)
```

```
## note: only 7 unique complexity parameters in default grid. Truncating the grid to 7 .
```

```
## Warning in train.default(x, y, weights = w, ...): The metric "Accuracy" was not  
## in the result set. ROC will be used instead.
```

```
#Accuracy:
```

```
rf_pred_test <- predict(rf, newdata =testing[,1:8], type = "raw")  
caret::confusionMatrix(rf_pred_test, testing$Transport_car)
```

```
## Confusion Matrix and Statistics  
##  
##           Reference  
## Prediction  No  Yes  
##           No 127   3  
##           Yes   0   9  
##  
##           Accuracy : 0.9784  
##           95% CI : (0.9382, 0.9955)  
##           No Information Rate : 0.9137  
##           P-Value [Acc > NIR] : 0.001664  
##  
##           Kappa : 0.8457  
##  
##           McNemar's Test P-Value : 0.248213  
##  
##           Sensitivity : 1.0000  
##           Specificity : 0.7500  
##           Pos Pred Value : 0.9769  
##           Neg Pred Value : 1.0000  
##           Prevalence : 0.9137  
##           Detection Rate : 0.9137  
##           Detection Prevalence : 0.9353  
##           Balanced Accuracy : 0.8750  
##  
##           'Positive' Class : No  
##
```

bagging :

```
bagging_model<-train(  
  Transport_car~.,  
  data = training,  
  method = "treebag",  
  nleaves=10,  
  ntrees=5,  
  trControl=Crul,  
  importance=TRUE  
)
```

```
## Warning in train.default(x, y, weights = w, ...): The metric "Accuracy" was not  
## in the result set. ROC will be used instead.
```

```
#Accuracy:
```

```
bagging_predictions_test <- predict(bagging_model, newdata = testing, type = "raw")  
caret::confusionMatrix(bagging_predictions_test, testing$Transport_car)
```

```
## Confusion Matrix and Statistics  
##  
##           Reference  
## Prediction  No Yes  
##           No 127  3  
##           Yes  0  9  
##  
##           Accuracy : 0.9784  
##           95% CI : (0.9382, 0.9955)  
##           No Information Rate : 0.9137  
##           P-Value [Acc > NIR] : 0.001664  
##  
##           Kappa : 0.8457  
##  
##           McNemar's Test P-Value : 0.248213  
##  
##           Sensitivity : 1.0000  
##           Specificity : 0.7500  
##           Pos Pred Value : 0.9769  
##           Neg Pred Value : 1.0000  
##           Prevalence : 0.9137  
##           Detection Rate : 0.9137  
##           Detection Prevalence : 0.9353  
##           Balanced Accuracy : 0.8750  
##  
##           'Positive' Class : No  
##
```

xgboost: (without SMOTE)

```
xgb.grid <- expand.grid(nrounds = 150,
                      eta = c(0.01),
                      max_depth = c(4,7),
                      gamma = 0,           #default=0
                      colsample_bytree = 1, #default=1
                      min_child_weight = 1, #default=1
                      subsample = 1       #default=1
)
xgb_model <- train(Transport_car~.,
                  data=training,
                  method="xgbTree",
                  trControl=Crul,
                  tuneGrid=xgb.grid,
                  verbose=T,
                  nthread = 2
)
```

```
## Warning in train.default(x, y, weights = w, ...): The metric "Accuracy" was not
## in the result set. ROC will be used instead.
```

#Accuracy:

```
xgb_predictions_test <- predict(xgb_model, newdata = testing, type = "raw")
confusionMatrix(xgb_predictions_test, testing$Transport_car)
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction  No Yes
##           No 127  1
##           Yes  0 11
##
##           Accuracy : 0.9928
##           95% CI : (0.9606, 0.9998)
##           No Information Rate : 0.9137
##           P-Value [Acc > NIR] : 5.011e-05
##
##           Kappa : 0.9526
##
##           Mcnemar's Test P-Value : 1
##
##           Sensitivity : 1.0000
##           Specificity : 0.9167
##           Pos Pred Value : 0.9922
##           Neg Pred Value : 1.0000
##           Prevalence : 0.9137
##           Detection Rate : 0.9137
##           Detection Prevalence : 0.9209
##           Balanced Accuracy : 0.9583
```

```
##
##      'Positive' Class : No
##
```

SMOTE:

```
table(training$Transport_car)
```

```
##
## No Yes
## 255  23
```

```
prop.table(table(training$Transport_car))
```

```
##
##      No      Yes
## 0.91726619 0.08273381
```

```
smote_train <- SMOTE(Transport_car ~ ., data = training,
                     perc.over = 3000,
                     perc.under = 300,
                     k = 5)
```

```
prop.table(table(smote_train$Transport_car))
```

```
##
##      No      Yes
## 0.7438017 0.2561983
```

```
table(smote_train$Transport_car)
```

```
##
## No  Yes
## 2070 713
```

xgboost: (with SMOTE)

```
xgb.grid <- expand.grid(nrounds = 150,
                      eta = c(0.01),
                      max_depth = c(4,7),
                      gamma = 0,                                     #default=0
                      colsample_bytree = 1,                         #default=1
                      min_child_weight = 1,                         #default=1
                      subsample = 1,                                 #default=1
                      )
smote_model <- train(Transport_car ~ .,
```

```

        data=smote_train,
        method="xgbTree",
        trControl=Crul,
        tuneGrid=xgb.grid,
        verbose=T,
        nthread = 2
    )

```

```

## Warning in train.default(x, y, weights = w, ...): The metric "Accuracy" was not
## in the result set. ROC will be used instead.

```

#Accuracy:

```

smote_predictions_test <- predict(smote_model, newdata = testing, type = "raw")
confusionMatrix(smote_predictions_test, testing$Transport_car)

```

```

## Confusion Matrix and Statistics
##
##           Reference
## Prediction  No Yes
##           No 127  1
##           Yes  0 11
##
##           Accuracy : 0.9928
##           95% CI : (0.9606, 0.9998)
##           No Information Rate : 0.9137
##           P-Value [Acc > NIR] : 5.011e-05
##
##           Kappa : 0.9526
##
##           Mcnemar's Test P-Value : 1
##
##           Sensitivity : 1.0000
##           Specificity : 0.9167
##           Pos Pred Value : 0.9922
##           Neg Pred Value : 1.0000
##           Prevalence : 0.9137
##           Detection Rate : 0.9137
##           Detection Prevalence : 0.9209
##           Balanced Accuracy : 0.9583
##
##           'Positive' Class : No
##

```

COMPARING MODELS

```

# Compare model performances using resample()
models_to_compare <- resamples(list(Logistic_Regression = glm,
                                     Navie_Bayes = naive_base,
                                     KNN = knn,

```

```

        bagging = bagging_model,
        Single_tree = rpart_model,
        smote=smote_model,
        Random_Forest = rf,
        xgboost=xgb_model
    ))

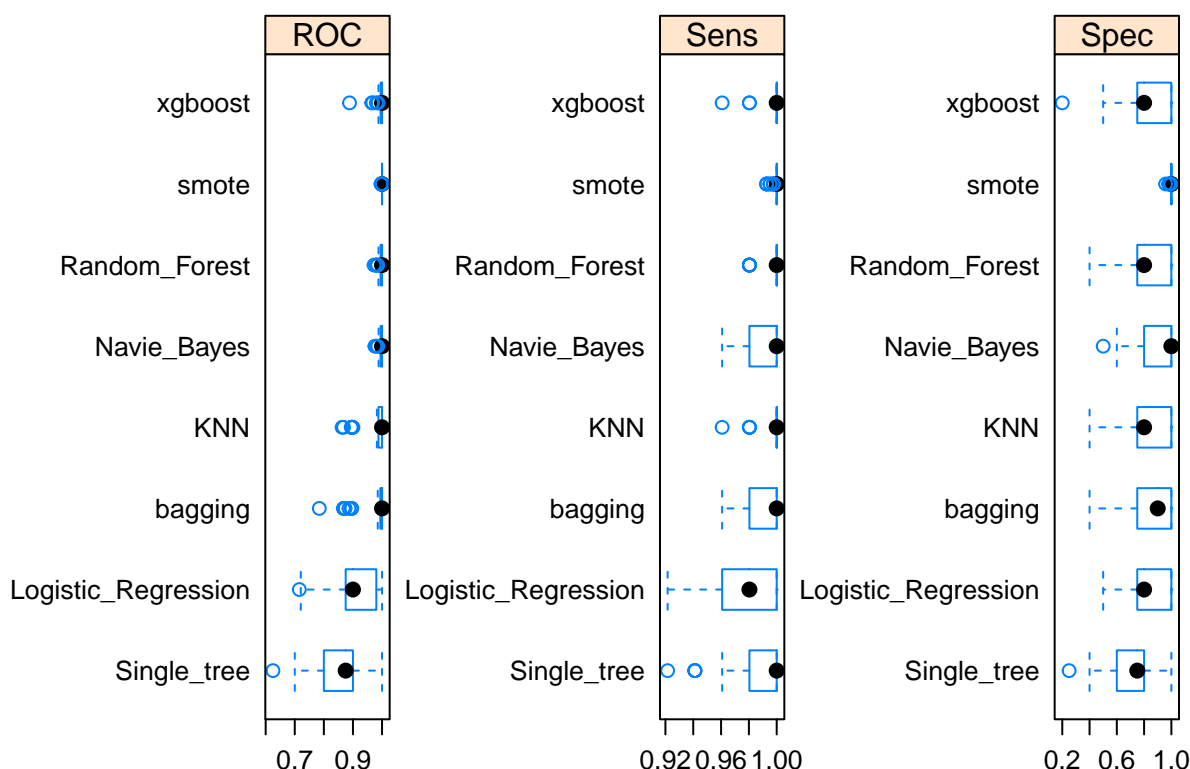
# Summary of the models performances
summary(models_to_compare)

##
## Call:
## summary.resamples(object = models_to_compare)
##
## Models: Logistic_Regression, Navie_Bayes, KNN, bagging, Single_tree, smote, Random_Forest, xgboost
## Number of resamples: 50
##
## ROC
##
##           Min.    1st Qu. Median      Mean   3rd Qu. Max. NA's
## Logistic_Regression 0.7156863 0.8750000  0.900 0.9175196 0.9803922    1    0
## Navie_Bayes         0.9754902 0.9950980  1.000 0.9963725 1.0000000    1    0
## KNN                 0.8627451 0.9883578  1.000 0.9793333 1.0000000    1    0
## bagging             0.7843137 0.9946078  1.000 0.9754608 1.0000000    1    0
## Single_tree         0.6250000 0.8000000  0.875 0.8657941 0.9000000    1    0
## smote               0.9963003 1.0000000  1.000 0.9999159 1.0000000    1    0
## Random_Forest       0.9725490 0.9950980  1.000 0.9960098 1.0000000    1    0
## xgboost             0.8882353 0.9950980  1.000 0.9941373 1.0000000    1    0
##
## Sens
##
##           Min.    1st Qu.   Median      Mean 3rd Qu. Max. NA's
## Logistic_Regression 0.9215686 0.9607843 0.9803922 0.9780392    1    1    0
## Navie_Bayes         0.9607843 0.9803922 1.0000000 0.9882353    1    1    0
## KNN                 0.9607843 1.0000000 1.0000000 0.9976471    1    1    0
## bagging             0.9607843 0.9852941 1.0000000 0.9941176    1    1    0
## Single_tree         0.9215686 0.9803922 1.0000000 0.9866667    1    1    0
## smote               0.9927536 1.0000000 1.0000000 0.9991787    1    1    0
## Random_Forest       0.9803922 1.0000000 1.0000000 0.9972549    1    1    0
## xgboost             0.9607843 1.0000000 1.0000000 0.9984314    1    1    0
##
## Spec
##
##           Min. 1st Qu. Median      Mean 3rd Qu. Max. NA's
## Logistic_Regression 0.5000000  0.75  0.80 0.8340000    1.0    1    0
## Navie_Bayes         0.5000000  0.80  1.00 0.8980000    1.0    1    0
## KNN                 0.4000000  0.75  0.80 0.8330000    1.0    1    0
## bagging             0.4000000  0.75  0.90 0.8590000    1.0    1    0
## Single_tree         0.2500000  0.60  0.75 0.7410000    0.8    1    0
## smote               0.9577465  1.00  1.00 0.9977534    1.0    1    0
## Random_Forest       0.4000000  0.75  0.80 0.8230000    1.0    1    0
## xgboost             0.2000000  0.75  0.80 0.8190000    1.0    1    0

```

Draw box plots to compare models

```
scales <- list(x=list(relation="free"), y=list(relation="free"))
bwplot(models_to_compare, scales=scales)
```



output: pdf_document

Summary:

1- we try to understand what transport employees prefers to commute to their office Car or other , so we upload and Data Preparation and split data in to two part Train,Test and we applied multiple “7” models with general parameters. 2- lets discuss the result : biased on the best accuracy : “bagging, xgboost, random forest, naive baise” had same acuuracy value : 99.28% , after that “smote, Logistic Regression” had Accuracy of : 98.56%, befor last “knn” with accuracy of : 97.84% , and last “Single decision tree” with :96.4%. 3- but when sort biased on ROC and Sensitivity and Specificity : witch are very important for choosing the model, “smote” is the best of all, and then “bagging,naive baise” . 4- the last disussion : since smote is the best on ROC and Sensitivity and Specificity and had Accuracy of : 98.56% witch is Good. I will go with smote model.